

Fast and deterministic system tests

Stevan Andjelkovic and Daniel Gustafsson

28th Oct, 2020

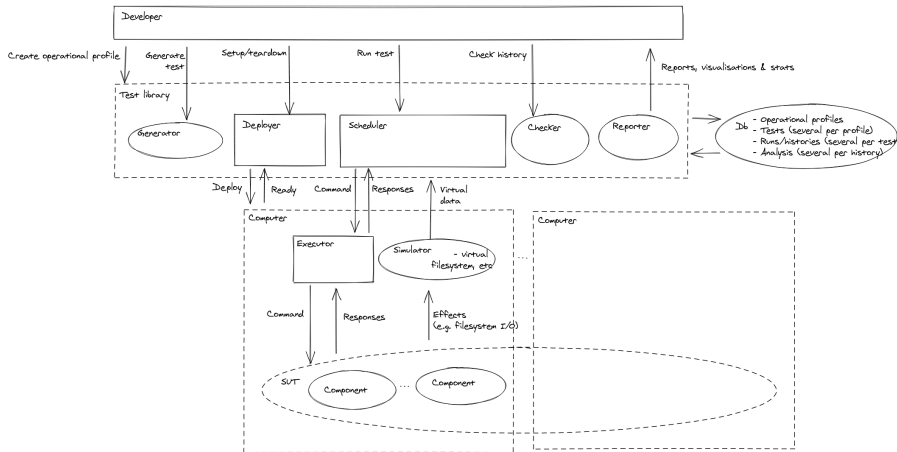
The problem

System tests, in general, are:

- ▶ Non-deterministic
 - ▶ Running the same test twice can yield different outcomes, esp. around fault-injection;
- ▶ Ill-specified or provide weak guarantees
 - ▶ What exactly have we shown if the tests pass?
- ▶ Ephemeral
 - ▶ Can't test performance over time;
 - ▶ Can't test upgrades, or backup and restore;
- ▶ Language specific
 - ▶ Test libraries/frameworks/tools are programming language specific, while the components of systems under test are written in different languages.

Parts of the solution

- ▶ Generator: generates random test cases;
- ▶ Scheduler: deterministically controls the network traffic during the test;
- ▶ Executor: receives messages from the scheduler and executes them against the system under test (SUT);
- ▶ Injector: figures out which faults to inject;
- ▶ Checker: analyses the output of a test case execution and determines if it was a success or not.



Solution for non-determinism

- ▶ SUT is assumed to be written on reactor form, i.e. given a message and some state, produce a set of messages;
- ▶ All messages get set via the Scheduler which randomly, but deterministically using a seed, determines the arrival order of the messages;
- ▶ Timeouts and retires are handled by explicit tick messages, that are also deterministically sent by the Scheduler.

Language agnostic solution

- ▶ In between the SUT and the Scheduler sits the Executor, whose job is to receive messages from the Scheduler via an http interface and pass them on to the SUT;
- ▶ The Executor is written in the same language as the SUT, so once it got the message via http it decodes the message from JSON into a datastructure in the native language and does a simple function call to the the SUT;
- ▶ Porting an Executor to a new programming language is simple, which means it's easy to test systems written using many languages.

Solution to ill-specified guarantees

- ▶ The Checker component uses Jepsen's state-of-the-art Elle checker, which provides precise models and guarantees;
- ▶ Lineage-driven fault injection is used to give guarantees in the presence faults;
- ▶ Operational profiles/usage models will later be used to drive test case generation, and guarantee system test coverage.

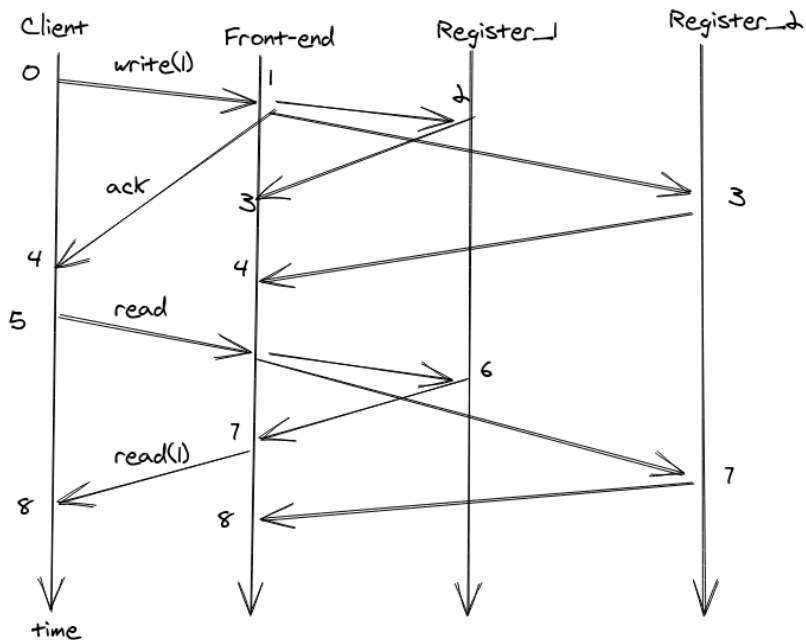
Solution to long-lived testing

- ▶ Every interaction that the developer can do, e.g. generation, execution, checking, can be done in isolation because the input and output comes and goes via a database;
- ▶ The above in combination with determinism means that we can replay an old test and bring the system to the state it was in at the end of a test, we can then extend the test can carry on from there.


Demo: the SUT

- ▶ The example SUT is a integer-valued shared/distributed register;
- ▶ Any number of clients can write or read an integer from the register;
- ▶ The register is replicated to try to achieve fault tolerance.

Demo: shared register v1, success



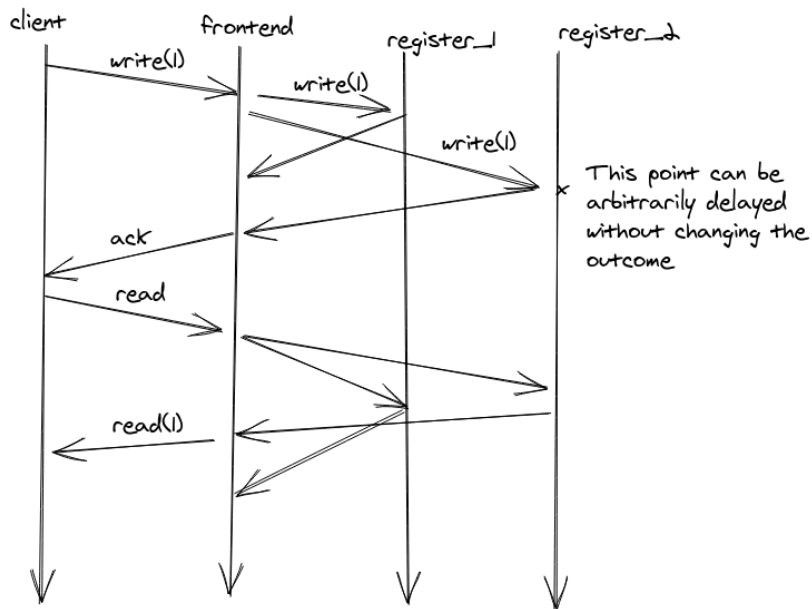
Demo: shared register v1, counterexample




Demo: the testsuite of the SUT

- ▶ Show the code of `detsys/sut/register_test.go`;
- ▶ `go test`;
- ▶ Ensure that we find the problem.

Example: shared register v2, success



Example: shared register v2, counterexample



Bonus: How does the Elle checker work?

Bonus: How does lineage-driven fault injection work?

Future work (next release)

- ▶ Regression tests
- ▶ Integration with Sean's work

Questions?